

example

January 23, 2024

```
[9]: import matplotlib.pyplot as plt
from ucimlrepo import fetch_ucirepo
import pandas as pd
import numpy as np
```

```
[10]: mushroom = fetch_ucirepo(id=73)
mushroom_features_df = pd.DataFrame.from_dict(mushroom['data']['features'])
mushroom_classes_df = pd.DataFrame.from_dict(mushroom['data']['targets'])
pd.concat([mushroom_classes_df, mushroom_features_df], axis=1)
```

```
[10]:    poisonous cap-shape cap-surface cap-color bruises odor gill-attachment \
0          p          x          s          n          t          p          f
1          e          x          s          y          t          a          f
2          e          b          s          w          t          l          f
3          p          x          y          w          t          p          f
4          e          x          s          g          f          n          f
...      ...      ...      ...      ...      ...      ...      ...
8119         e          k          s          n          f          n          a
8120         e          x          s          n          f          n          a
8121         e          f          s          n          f          n          a
8122         p          k          y          n          f          y          f
8123         e          x          s          n          f          n          a
```

```
    gill-spacing gill-size gill-color ... stalk-surface-below-ring \
0              c          n          k ...                          s
1              c          b          k ...                          s
2              c          b          n ...                          s
3              c          n          n ...                          s
4              w          b          k ...                          s
...      ...      ...      ...      ...      ...
8119         c          b          y ...                          s
8120         c          b          y ...                          s
8121         c          b          n ...                          s
8122         c          n          b ...                          k
8123         c          b          y ...                          s
```

```
    stalk-color-above-ring stalk-color-below-ring veil-type veil-color \
0              w          w          p          w
```

1		w		w	p	w
2		w		w	p	w
3		w		w	p	w
4		w		w	p	w
...	
8119		o		o	p	o
8120		o		o	p	n
8121		o		o	p	o
8122		w		w	p	w
8123		o		o	p	o

	ring-number	ring-type	spore-print-color	population	habitat	
0		o	p	k	s	u
1		o	p	n	n	g
2		o	p	n	n	m
3		o	p	k	s	u
4		o	e	n	a	g
...	
8119		o	p	b	c	l
8120		o	p	b	v	l
8121		o	p	b	c	l
8122		o	e	w	v	l
8123		o	p	o	c	l

[8124 rows x 23 columns]

```
[11]: breast_cancer = fetch_ucirepo(id=17)
breast_cancer_features_df = pd.DataFrame.
    ↪from_dict(breast_cancer['data']['features'])
breast_cancer_classes_df = pd.DataFrame.
    ↪from_dict(breast_cancer['data']['targets'])
pd.concat([breast_cancer_classes_df, breast_cancer_features_df], axis=1)
```

```
[11]:   Diagnosis  radius1  texture1  perimeter1  area1  smoothness1  \
0         M    17.99    10.38    122.80  1001.0    0.11840
1         M    20.57    17.77    132.90  1326.0    0.08474
2         M    19.69    21.25    130.00  1203.0    0.10960
3         M    11.42    20.38     77.58   386.1    0.14250
4         M    20.29    14.34    135.10  1297.0    0.10030
..      ...      ...      ...      ...      ...      ...
564        M    21.56    22.39    142.00  1479.0    0.11100
565        M    20.13    28.25    131.20  1261.0    0.09780
566        M    16.60    28.08    108.30   858.1    0.08455
567        M    20.60    29.33    140.10  1265.0    0.11780
568        B     7.76    24.54     47.92   181.0    0.05263

compactness1  concavity1  concave_points1  symmetry1  ...  radius3  \
```

0	0.27760	0.30010	0.14710	0.2419	...	25.380
1	0.07864	0.08690	0.07017	0.1812	...	24.990
2	0.15990	0.19740	0.12790	0.2069	...	23.570
3	0.28390	0.24140	0.10520	0.2597	...	14.910
4	0.13280	0.19800	0.10430	0.1809	...	22.540
..
564	0.11590	0.24390	0.13890	0.1726	...	25.450
565	0.10340	0.14400	0.09791	0.1752	...	23.690
566	0.10230	0.09251	0.05302	0.1590	...	18.980
567	0.27700	0.35140	0.15200	0.2397	...	25.740
568	0.04362	0.00000	0.00000	0.1587	...	9.456

	texture3	perimeter3	area3	smoothness3	compactness3	concavity3	\
0	17.33	184.60	2019.0	0.16220	0.66560	0.7119	
1	23.41	158.80	1956.0	0.12380	0.18660	0.2416	
2	25.53	152.50	1709.0	0.14440	0.42450	0.4504	
3	26.50	98.87	567.7	0.20980	0.86630	0.6869	
4	16.67	152.20	1575.0	0.13740	0.20500	0.4000	
..	
564	26.40	166.10	2027.0	0.14100	0.21130	0.4107	
565	38.25	155.00	1731.0	0.11660	0.19220	0.3215	
566	34.12	126.70	1124.0	0.11390	0.30940	0.3403	
567	39.42	184.60	1821.0	0.16500	0.86810	0.9387	
568	30.37	59.16	268.6	0.08996	0.06444	0.0000	

	concave_points3	symmetry3	fractal_dimension3
0	0.2654	0.4601	0.11890
1	0.1860	0.2750	0.08902
2	0.2430	0.3613	0.08758
3	0.2575	0.6638	0.17300
4	0.1625	0.2364	0.07678
..
564	0.2216	0.2060	0.07115
565	0.1628	0.2572	0.06637
566	0.1418	0.2218	0.07820
567	0.2650	0.4087	0.12400
568	0.0000	0.2871	0.07039

[569 rows x 31 columns]

```
[12]: from sklearn.preprocessing import LabelEncoder, LabelBinarizer
from sklearn.datasets import fetch_openml

titanic = fetch_openml(name='titanic', version=1)
titanic_df = titanic['frame']
titanic_df['boat'] = titanic_df['boat'].fillna(value=0)
titanic_df['body'] = titanic_df['body'].fillna(value=0)
```

```
titanic_df['body'] = titanic_df['body'].fillna(value=0)
```

```
titanic_df
```

```
[12]:
```

	pclass	survived	name \
0	1	1	Allen, Miss. Elisabeth Walton
1	1	1	Allison, Master. Hudson Trevor
2	1	0	Allison, Miss. Helen Loraine
3	1	0	Allison, Mr. Hudson Joshua Creighton
4	1	0	Allison, Mrs. Hudson J C (Bessie Waldo Daniels)
...
1304	3	0	Zabour, Miss. Hileni
1305	3	0	Zabour, Miss. Thamine
1306	3	0	Zakarian, Mr. Mapriededer
1307	3	0	Zakarian, Mr. Ortin
1308	3	0	Zimmerman, Mr. Leo

	sex	age	sibsp	parch	ticket	fare	cabin	embarked	boat \
0	female	29.0000	0	0	24160	211.3375	B5	S	2
1	male	0.9167	1	2	113781	151.5500	C22 C26	S	11
2	female	2.0000	1	2	113781	151.5500	C22 C26	S	0
3	male	30.0000	1	2	113781	151.5500	C22 C26	S	0
4	female	25.0000	1	2	113781	151.5500	C22 C26	S	0
...
1304	female	14.5000	1	0	2665	14.4542	NaN	C	0
1305	female	NaN	1	0	2665	14.4542	NaN	C	0
1306	male	26.5000	0	0	2656	7.2250	NaN	C	0
1307	male	27.0000	0	0	2670	7.2250	NaN	C	0
1308	male	29.0000	0	0	315082	7.8750	NaN	S	0

	body	home.dest
0	0.0	St Louis, MO
1	0.0	Montreal, PQ / Chesterville, ON
2	0.0	Montreal, PQ / Chesterville, ON
3	135.0	Montreal, PQ / Chesterville, ON
4	0.0	Montreal, PQ / Chesterville, ON
...
1304	328.0	NaN
1305	0.0	NaN
1306	304.0	NaN
1307	0.0	NaN
1308	0.0	NaN

```
[1309 rows x 14 columns]
```

```
[13]: titanic_df_encoded = titanic_df

titanic_df['boat'] = titanic_df['boat'].astype(str)
titanic_df['boat'] = titanic_df['boat'].astype(str)

encoded_features = ['sex', 'cabin', 'ticket', 'home.dest', 'embarked', 'boat']
label_encoders = {feature: LabelEncoder() for feature in encoded_features}
for feature, encoder in label_encoders.items():
    titanic_df_encoded[feature] = encoder.fit_transform(titanic_df[feature])

titanic_lb_survived = LabelBinarizer()
titanic_df_encoded['survived'] = titanic_lb_survived.
    ↪fit_transform(titanic_df['survived'])

titanic_df_encoded
```

```
[13]:
```

	pclass	survived	name	sex	\
0	1	1	Allen, Miss. Elisabeth Walton	0	
1	1	1	Allison, Master. Hudson Trevor	1	
2	1	0	Allison, Miss. Helen Loraine	0	
3	1	0	Allison, Mr. Hudson Joshua Creighton	1	
4	1	0	Allison, Mrs. Hudson J C (Bessie Waldo Daniels)	0	
...
1304	3	0	Zabour, Miss. Hileni	0	
1305	3	0	Zabour, Miss. Thamine	0	
1306	3	0	Zakarian, Mr. Mapriededer	1	
1307	3	0	Zakarian, Mr. Ortin	1	
1308	3	0	Zimmerman, Mr. Leo	1	

	age	sibsp	parch	ticket	fare	cabin	embarked	boat	body	\
0	29.0000	0	0	187	211.3375	43	2	12	0.0	
1	0.9167	1	2	49	151.5500	79	2	3	0.0	
2	2.0000	1	2	49	151.5500	79	2	0	0.0	
3	30.0000	1	2	49	151.5500	79	2	0	135.0	
4	25.0000	1	2	49	151.5500	79	2	0	0.0	
...
1304	14.5000	1	0	259	14.4542	186	0	0	328.0	
1305	NaN	1	0	259	14.4542	186	0	0	0.0	
1306	26.5000	0	0	250	7.2250	186	0	0	304.0	
1307	27.0000	0	0	264	7.2250	186	0	0	0.0	
1308	29.0000	0	0	346	7.8750	186	2	0	0.0	

	home.dest
0	308
1	230
2	230
3	230

```

4          230
...
1304       369
1305       369
1306       369
1307       369
1308       369

```

[1309 rows x 14 columns]

0.0.1 Titanic example

```

[14]: def evaluate(clf, X, y, scoring=None) -> dict:
        scoring = scoring or ['accuracy', 'precision', 'recall', 'roc_auc', 'f1']

        cv = ShuffleSplit(n_splits=5, test_size=0.2, random_state=0)
        _scores = cross_validate(clf, X, y, cv=cv, scoring=scoring, n_jobs=8,
        ↪return_train_score=True)
        return {name: score.mean() for name, score in _scores.items()}

def evaluate_return_estimator(clf, X, y, scoring=None) -> tuple[dict, list]:
        scoring = scoring or ['accuracy', 'precision', 'recall', 'roc_auc', 'f1']

        cv = ShuffleSplit(n_splits=5, test_size=0.2, random_state=0)
        _scores = cross_validate(clf, X, y, cv=cv, scoring=scoring, n_jobs=8,
        ↪return_train_score=True,
                                return_estimator=True)
        return {name: score.mean() for name, score in _scores.items() if name !=
        ↪'estimator'}, _scores['estimator']

```

```

[15]: from sklearn.ensemble import RandomForestClassifier

```

```

[16]: from sklearn.model_selection import cross_validate, ShuffleSplit

N_ESTIMATORS = 200

clf = RandomForestClassifier(n_estimators=N_ESTIMATORS)

X = titanic_df_encoded.drop(['survived', 'name'], axis='columns')
y = titanic_df['survived']

evaluate(clf, X, y, scoring=['accuracy', 'precision', 'recall', 'roc_auc',
        ↪'f1'])

```

```

[16]: {'fit_time': 0.31610894203186035,
        'score_time': 0.030126237869262697,

```

```

'test_accuracy': 0.9702290076335878,
'train_accuracy': 0.9793696275071634,
'test_precision': 0.9787358059706378,
'train_precision': 0.9821955063295645,
'test_recall': 0.9408521106374348,
'train_recall': 0.9635912785219982,
'test_roc_auc': 0.9932495116763438,
'train_roc_auc': 0.9982890332080114,
'test_f1': 0.9594007649305427,
'train_f1': 0.9728011290549146}

```

0.0.2 Breast cancer example

```

[17]: clf = RandomForestClassifier(n_estimators=N_ESTIMATORS)

X = breast_cancer_features_df
breast_cancer_lb_diagnosis = LabelBinarizer()

X = breast_cancer_features_df.to_numpy()
y = np.squeeze(breast_cancer_lb_diagnosis.
    ↪fit_transform(breast_cancer_classes_df))

evaluate(clf, X, y, scoring=['accuracy', 'precision', 'recall', 'roc_auc',
    ↪'f1'])

```

```

[17]: {'fit_time': 0.3469456672668457,
'score_time': 0.02559514045715332,
'test_accuracy': 0.9631578947368421,
'train_accuracy': 1.0,
'test_precision': 0.9698643410852712,
'train_precision': 1.0,
'test_recall': 0.936071449815499,
'train_recall': 1.0,
'test_roc_auc': 0.9909137176413665,
'train_roc_auc': 1.0,
'test_f1': 0.9516461901101048,
'train_f1': 1.0}

```

0.0.3 Mushroom example

```

[63]: clf = RandomForestClassifier(n_estimators=N_ESTIMATORS)

mushroom_lb_poisonous = LabelBinarizer()

mushroom_features_df_encoded = mushroom_features_df

encoded_features = mushroom_features_df.columns

```

```

label_encoders = {feature: LabelEncoder() for feature in encoded_features}
for feature, encoder in label_encoders.items():
    mushroom_features_df[feature] = encoder.
    ↪fit_transform(mushroom_features_df[feature])

X = mushroom_features_df_encoded
y = np.squeeze(mushroom_lb_poisonous.fit_transform(mushroom_classes_df))

evaluate(clf, X, y, scoring=['accuracy', 'precision', 'recall', 'roc_auc',
    ↪'f1'])

```

```

[63]: {'fit_time': 0.4403873920440674,
      'score_time': 0.046620416641235354,
      'test_accuracy': 1.0,
      'train_accuracy': 1.0,
      'test_precision': 1.0,
      'train_precision': 1.0,
      'test_recall': 1.0,
      'train_recall': 1.0,
      'test_roc_auc': 1.0,
      'train_roc_auc': 1.0,
      'test_f1': 1.0,
      'train_f1': 1.0}

```

1 Boosted RF

Plain ol' AdaBoost

```

[69]: from sklearn.ensemble import AdaBoostClassifier

clf = AdaBoostClassifier(n_estimators=N_ESTIMATORS, algorithm='SAMME')

X = breast_cancer_features_df
breast_cancer_lb_diagnosis = LabelBinarizer()

y = np.squeeze(breast_cancer_lb_diagnosis.
    ↪fit_transform(breast_cancer_classes_df))

evaluate(clf, X, y, scoring=['accuracy', 'precision', 'recall', 'roc_auc',
    ↪'f1'])

```

```

[69]: {'fit_time': 0.5341531753540039,
      'score_time': 0.052882146835327146,
      'test_accuracy': 0.9719298245614034,
      'train_accuracy': 1.0,
      'test_precision': 0.9769461694209735,
      'train_precision': 1.0,

```



```

'test_recall': 0.9497502722903775,
'train_recall': 1.0,
'test_roc_auc': 0.9966182602092898,
'train_roc_auc': 1.0,
'test_f1': 0.9629576928742024,
'train_f1': 1.0}

```

“Boosted RF”

```

[20]: from IPython.display import display
      from sklearn.tree import DecisionTreeClassifier

      BOOST_RF_N_ESTIMATORS = 50

      clf_ref = AdaBoostClassifier(n_estimators=BOOST_RF_N_ESTIMATORS,
                                   ↪algorithm='SAMME')

      class BoostedRFClassifier(AdaBoostClassifier):
          def __init__(
              self,
              *,
              estimator=DecisionTreeClassifier(),
              n_estimators=BOOST_RF_N_ESTIMATORS,
              learning_rate=1.0,
              random_state=None,
              algorithm='SAMME',
              _disable_weights=False
          ):
              super().__init__(
                  estimator=estimator,
                  n_estimators=n_estimators,
                  learning_rate=learning_rate,
                  random_state=random_state,
                  algorithm=algorithm
              )

              self._disable_weights = _disable_weights
              self._estimator_weights = np.zeros(self.n_estimators, dtype=np.float64)

              @property
              def estimator_weights(self):
                  return np.ones(self.n_estimators, dtype=np.float64) if self.
                  ↪_disable_weights else self._estimator_weights

              @estimator_weights.setter
              def estimator_weights(self, v):

```

```

        if not self._disable_weights:
            self._estimator_weights = v

clf = BoostedRFClassifier(
    estimator=DecisionTreeClassifier(max_depth=5),
    n_estimators=BOOST_RF_N_ESTIMATORS,
    algorithm='SAMME',
    _disable_weights=True
)

X = breast_cancer_features_df
breast_cancer_lb_diagnosis = LabelBinarizer()

y = np.squeeze(breast_cancer_lb_diagnosis.
    ↪fit_transform(breast_cancer_classes_df))

res_ref = evaluate(clf_ref, X, y, scoring=['accuracy', 'precision', 'recall',
    ↪'roc_auc', 'f1'])
res, est = evaluate_return_estimator(clf, X, y, scoring=['accuracy',
    ↪'precision', 'recall', 'roc_auc', 'f1'])

display(
    pd.concat([
        pd.DataFrame.from_dict(res_ref, orient='index')[0].rename('AdaBoost'),
        pd.DataFrame.from_dict(res, orient='index')[0].rename('Boosted RF')
    ], axis=1)
)

```

	AdaBoost	Boosted RF
fit_time	0.144744	0.131094
score_time	0.020160	0.011134
test_accuracy	0.973684	0.961404
train_accuracy	1.000000	1.000000
test_precision	0.982013	0.972925
train_precision	1.000000	1.000000
test_recall	0.950175	0.927043
train_recall	1.000000	1.000000
test_roc_auc	0.994809	0.993863
train_roc_auc	1.000000	1.000000
test_f1	0.965594	0.948692
train_f1	1.000000	1.000000

```

[21]: tree_depths = list(range(1, 33))
measures = ['accuracy', 'precision', 'recall', 'roc_auc', 'f1']
res = [evaluate(BoostedRFClassifier(
    estimator=DecisionTreeClassifier(max_depth=depth),

```

```

        n_estimators=BOOST_RF_N_ESTIMATORS,
        algorithm='SAMME',
        _disable_weights=False
    ), X, y, scoring=measures) for depth in tree_depths]

```

```

[22]: res_no_weights = [evaluate(BoostedRFClassifier(
    estimator=DecisionTreeClassifier(max_depth=depth),
    n_estimators=BOOST_RF_N_ESTIMATORS,
    algorithm='SAMME',
    _disable_weights=True
), X, y, scoring=measures) for depth in tree_depths]

```

```

[23]: res_ref = [evaluate(RandomForestClassifier(
    max_depth=depth,
    n_estimators=BOOST_RF_N_ESTIMATORS,
    n_jobs=8
), X, y, scoring=measures) for depth in tree_depths]

```

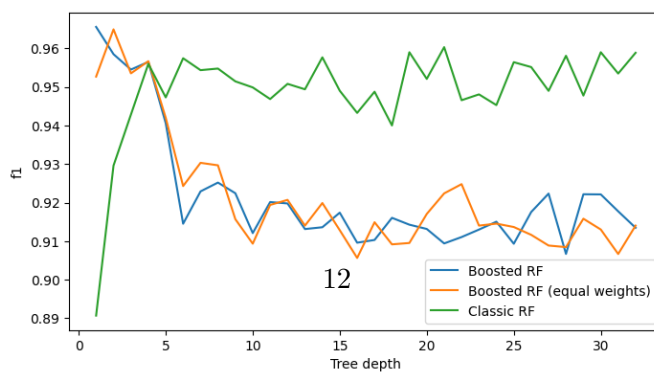


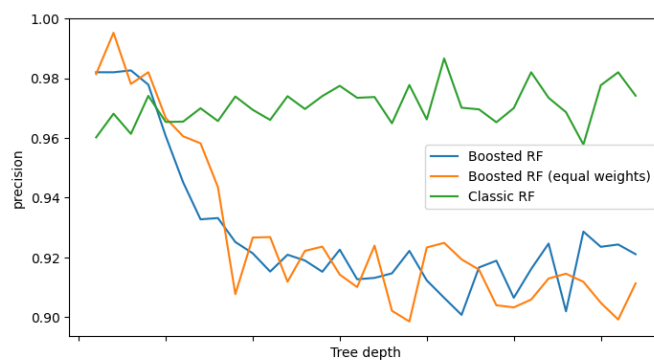

```

[24]: from matplotlib import pyplot as plt

fig, axes = plt.subplots(5, 1, sharex=True)
fig.set_size_inches(8, 25)
for measure, ax in zip(measures, axes):
    _y = [r[f"test_{measure}"] for r in res]
    _y_ref = [r[f"test_{measure}"] for r in res_ref]
    _y_no_weights = [r[f"test_{measure}"] for r in res_no_weights]
    ax.plot(tree_depths, _y, label="Boosted RF")
    ax.plot(tree_depths, _y_no_weights, label="Boosted RF (equal weights)")
    ax.plot(tree_depths, _y_ref, label="Classic RF")
    ax.set_xlabel("Tree depth")
    ax.set_ylabel(measure)
    ax.legend()

plt.show()

```



Stacking

```
[25]: from sklearn.ensemble import StackingClassifier

clf = StackingClassifier(estimators=[(f"tree{i}", RandomForestClassifier()) for
    ↳ i in range(16)],
                        final_estimator=RandomForestClassifier())

X = breast_cancer_features_df
breast_cancer_lb_diagnosis = LabelBinarizer()

y = np.squeeze(breast_cancer_lb_diagnosis.
    ↳ fit_transform(breast_cancer_classes_df))

evaluate(clf, X, y, scoring=['accuracy', 'precision', 'recall', 'roc_auc',
    ↳ 'f1'])
```

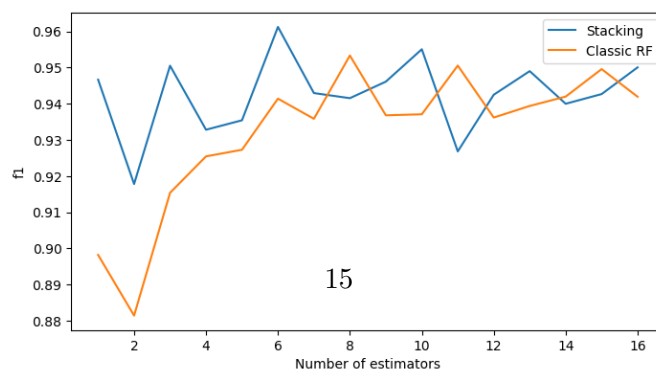


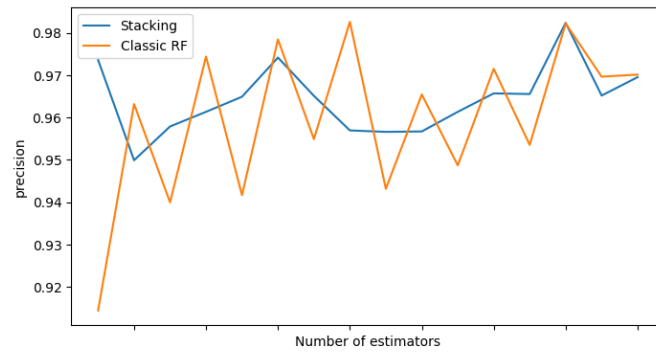

```
[25]: {'fit_time': 15.19865641593933,
      'score_time': 0.1616374969482422,
      'test_accuracy': 0.956140350877193,
      'train_accuracy': 0.9982417582417582,
      'test_precision': 0.969563887570532,
      'train_precision': 1.0,
      'test_recall': 0.9175271931425899,
      'train_recall': 0.9952438782242208,
      'test_roc_auc': 0.9894386455594674,
      'train_roc_auc': 0.9999958202716822,
      'test_f1': 0.9414258270465167,
      'train_f1': 0.9976113970689516}
```

```
[26]: num_estimators = list(range(1, 17))
measures = ['accuracy', 'precision', 'recall', 'roc_auc', 'f1']
res = [evaluate(StackingClassifier(estimators=[(f"tree{i}",
    ↳ RandomForestClassifier()) for i in range(n)],
                                final_estimator=RandomForestClassifier()),
    ↳ X, y, scoring=measures) for n in
    num_estimators]
```

```
[28]: res_ref = [evaluate(RandomForestClassifier(
    max_depth=None,
    n_estimators=num_estimator,
    n_jobs=8
), X, y, scoring=measures) for num_estimator in num_estimators]
```

```
[36]: fig, axes = plt.subplots(5, 1, sharex=True)
fig.set_size_inches(8, 25)
for measure, ax in zip(measures, axes):
    _y = [r[f"test_{measure}"] for r in res]
    _y_ref = [r[f"test_{measure}"] for r in res_ref]
    ax.plot(num_estimators, _y, label="Stacking")
    ax.plot(num_estimators, _y_ref, label="Classic RF")
    ax.set_xlabel("Number of estimators")
    ax.set_ylabel(measure)
    ax.legend()

plt.show()
```



```
[ ]: from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import LinearSVC
from sklearn.linear_model import LogisticRegression
```

```
clf = StackingClassifier(estimators=[
    ("nb", GaussianNB()),
    ("tree", DecisionTreeClassifier()),
    ("rf", RandomForestClassifier()),
    ("knn", KNeighborsClassifier()),
    ("svm", LinearSVC()),
    ("lr", LogisticRegression())
], final_estimator=RandomForestClassifier())
```

```
X = breast_cancer_features_df
breast_cancer_lb_diagnosis = LabelBinarizer()
```

```
y = np.squeeze(breast_cancer_lb_diagnosis.
    ↳fit_transform(breast_cancer_classes_df))
```

```
evaluate(clf, X, y, scoring=['accuracy', 'precision', 'recall', 'roc_auc',
    ↳'f1'])
```

```
[51]: res_single = evaluate(clf, X, y, scoring=['accuracy', 'precision', 'recall',
    ↳'roc_auc', 'f1'])
res_variety = [res_single] * len(num_estimators)
```

```
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`
will change from `True` to `auto` in 1.5. Set the value of `dual` explicitly
to suppress the warning.
```

```
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`
will change from `True` to `auto` in 1.5. Set the value of `dual` explicitly
to suppress the warning.
```

```
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
```

```
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
```



```
warnings.warn(  
/media/data/coding/uma/.venv/lib/python3.11/site-  
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  
/media/data/coding/uma/.venv/lib/python3.11/site-  
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`  
will change from `True` to `'auto'` in 1.5. Set the value of `dual` explicitly  
to suppress the warning.  
warnings.warn(  
/media/data/coding/uma/.venv/lib/python3.11/site-  
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  
/media/data/coding/uma/.venv/lib/python3.11/site-  
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to  
converge, increase the number of iterations.  
warnings.warn(  
/media/data/coding/uma/.venv/lib/python3.11/site-  
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`  
will change from `True` to `'auto'` in 1.5. Set the value of `dual` explicitly  
to suppress the warning.  
warnings.warn(  
/media/data/coding/uma/.venv/lib/python3.11/site-  
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`  
will change from `True` to `'auto'` in 1.5. Set the value of `dual` explicitly  
to suppress the warning.  
warnings.warn(  
/media/data/coding/uma/.venv/lib/python3.11/site-  
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`
will change from `True` to `'auto'` in 1.5. Set the value of `dual` explicitly
to suppress the warning.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`
will change from `True` to `'auto'` in 1.5. Set the value of `dual` explicitly
to suppress the warning.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`
will change from `True` to `'auto'` in 1.5. Set the value of `dual` explicitly
```

to suppress the warning.

```
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`
will change from `True` to `auto` in 1.5. Set the value of `dual` explicitly
to suppress the warning.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`
will change from `True` to `auto` in 1.5. Set the value of `dual` explicitly
to suppress the warning.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`
will change from `True` to `auto` in 1.5. Set the value of `dual` explicitly
to suppress the warning.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`
will change from `True` to `auto` in 1.5. Set the value of `dual` explicitly
to suppress the warning.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`
will change from `True` to `auto` in 1.5. Set the value of `dual` explicitly
to suppress the warning.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
```



```

warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`
will change from `True` to `'auto'` in 1.5. Set the value of `dual` explicitly
to suppress the warning.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`
will change from `True` to `'auto'` in 1.5. Set the value of `dual` explicitly
to suppress the warning.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`
will change from `True` to `'auto'` in 1.5. Set the value of `dual` explicitly
to suppress the warning.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`
will change from `True` to `'auto'` in 1.5. Set the value of `dual` explicitly
to suppress the warning.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to

```

```

converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`
will change from `True` to `auto` in 1.5. Set the value of `dual` explicitly
to suppress the warning.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`
will change from `True` to `auto` in 1.5. Set the value of `dual` explicitly
to suppress the warning.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed

```

```
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`
will change from `True` to `auto` in 1.5. Set the value of `dual` explicitly
to suppress the warning.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`
will change from `True` to `'auto'` in 1.5. Set the value of `dual` explicitly
to suppress the warning.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`
will change from `True` to `'auto'` in 1.5. Set the value of `dual` explicitly
to suppress the warning.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```


Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual`
will change from `True` to `auto` in 1.5. Set the value of `dual` explicitly
to suppress the warning.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
/media/data/coding/uma/.venv/lib/python3.11/site-  
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  
/media/data/coding/uma/.venv/lib/python3.11/site-  
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  
/media/data/coding/uma/.venv/lib/python3.11/site-  
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  
/media/data/coding/uma/.venv/lib/python3.11/site-  
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  
/media/data/coding/uma/.venv/lib/python3.11/site-  
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

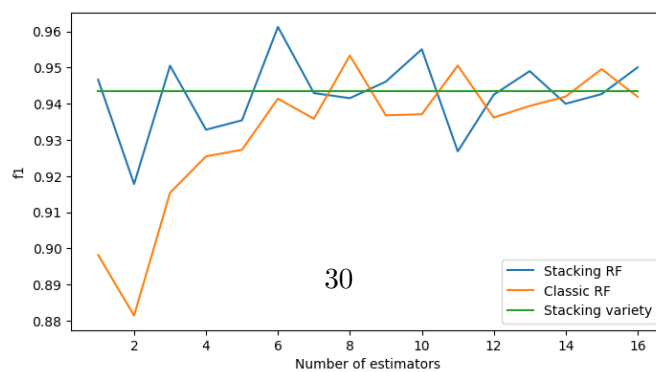




Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
[54]: fig, axes = plt.subplots(5, 1, sharex=True)
fig.set_size_inches(8, 25)
for measure, ax in zip(measures, axes):
    _y_rf = [r[f"test_{measure}"] for r in res]
    _y_variety = [r[f"test_{measure}"] for r in res_variety]
    _y_ref = [r[f"test_{measure}"] for r in res_ref]
    ax.plot(num_estimators, _y_rf, label="Stacking RF")
    ax.plot(num_estimators, _y_ref, label="Classic RF")
    ax.plot(num_estimators, _y_variety, label="Stacking variety")
    ax.set_xlabel("Number of estimators")
    ax.set_ylabel(measure)
    ax.legend()

plt.show()
```



1.1 Iterative RF

```
[104]: from sklearn.metrics import confusion_matrix
from sklearn.base import BaseEstimator

class IterativeRandomForrest(BaseEstimator):
    def __init__(
        self,
        *,
        iter_count=2,
        n_estimators=2,
        max_depth=None,
    ):

        self.iter_count = 2
        self.n_estimators = n_estimators
        self.max_depth = max_depth
        self._fitted_estimator = None

    def _get_estimator(self, weights):
        return RandomForestClassifier(
            n_estimators=self.n_estimators,
            max_depth=self.max_depth,
            n_jobs=8,
            class_weight=weights
        )

    def _get_accuracy_per_class(self, y, y_pred):
        c_matrix = confusion_matrix(y, y_pred)
        return c_matrix.diagonal() / c_matrix.sum(axis=1)

    def _get_weights(self, class_accuracies, y):
        unique_classes = np.unique(y)
        weights = {}
        for class_idx, class_accuracy in enumerate(class_accuracies):
            weights[unique_classes[class_idx]] = 1 - class_accuracy
        return weights

    def _get_initial_weights(self, y):
        unique_classes = np.unique(y)
        weights = {}
        for class_name in unique_classes:
            weights[class_name] = 1
        return weights
```

```

def fit(self, X, y):
    weights = self._get_initial_weights(y)

    for iteration_idx in range(self.iter_count):
        estimator = self._get_estimator(weights)
        estimator.fit(X, y)
        y_pred = estimator.predict(X)
        class_accuracies = self._get_accuracy_per_class(y, y_pred)
        weights = self._get_weights(class_accuracies, y)
        self._fitted_estimator = estimator

def predict(self, X):
    return self._fitted_estimator.predict(X)

def decision_function(self, X):
    # called for roc_auc_score
    raise NotImplementedError("This IterativeRandomForrest instance is not_
↳fitted yet")

def predict_proba(self, X):
    # called for roc_auc_score
    raise NotImplementedError("This IterativeRandomForrest instance is not_
↳fitted yet")

clf = IterativeRandomForrest(iter_count=2, n_estimators=2, max_depth=2)
clf.fit(X, y)
clf.predict(X)
clf.get_params()
evaluate(clf, X, y, scoring=['accuracy', 'precision', 'recall', 'f1'])

```

```

[104]: {'fit_time': 0.06289176940917969,
'score_time': 0.01744365692138672,
'test_accuracy': 0.9087719298245615,
'train_accuracy': 0.9261538461538462,
'test_precision': 0.8598644705347851,
'train_precision': 0.8698541768515575,
'test_recall': 0.9231330930751069,
'train_recall': 0.9460505956384206,
'test_f1': 0.8882752413680249,
'train_f1': 0.9043734538199694}

```

```

[105]: NUM_ESTIMATORS = 2
MAX_DEPTH = 2
iteration_counts = list(range(1, 9))
measures = ['accuracy', 'precision', 'recall', 'f1']

```



```

res = [evaluate(IterativeRandomForrest(n_estimators=NUM_ESTIMATORS,
    ↪max_depth=MAX_DEPTH, iter_count=iter_count), X, y, scoring=measures) for
    ↪iter_count in iteration_counts]

res_ref_single = evaluate(RandomForestClassifier(
    max_depth=MAX_DEPTH,
    n_estimators=NUM_ESTIMATORS,
    n_jobs=8
), X, y, scoring=measures)
res_ref = [res_ref_single] * len(iteration_counts)


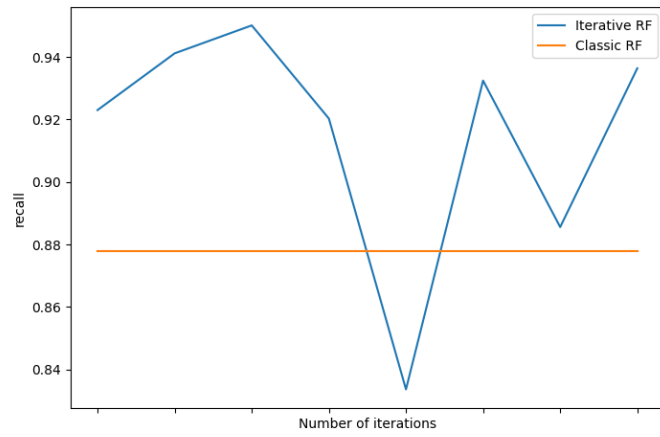
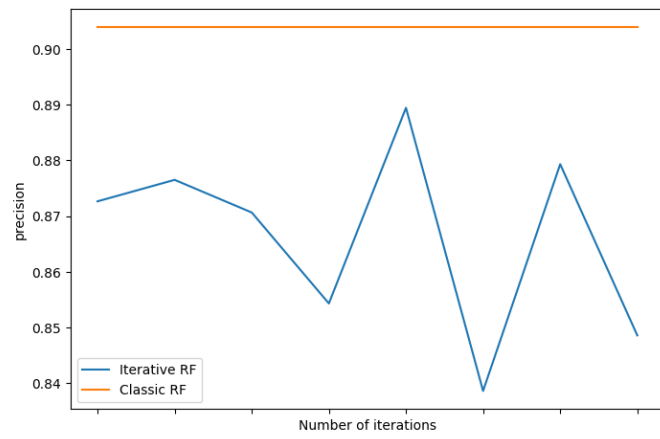

```

```

[106]: fig, axes = plt.subplots(4, 1, sharex=True)
fig.set_size_inches(8, 25)
for measure, ax in zip(measures, axes):
    _y_iter = [r[f"test_{measure}"] for r in res]
    _y_ref = [r[f"test_{measure}"] for r in res_ref]
    ax.plot(iteration_counts, _y_iter, label="Iterative RF")
    ax.plot(iteration_counts, _y_ref, label="Classic RF")
    ax.set_xlabel("Number of iterations")
    ax.set_ylabel(measure)
    ax.legend()

plt.show()

```



[]:

[107]: `from ucimlrepo import fetch_ucirepo`

```
# fetch dataset
statlog_german_credit_data = fetch_ucirepo(id=144)

# data (as pandas dataframes)
X = statlog_german_credit_data.data.features
y = statlog_german_credit_data.data.targets
```

```
{'uci_id': 144, 'name': 'Statlog (German Credit Data)', 'repository_url':
'https://archive.ics.uci.edu/dataset/144/statlog+german+credit+data',
'data_url': 'https://archive.ics.uci.edu/static/public/144/data.csv',
'abstract': 'This dataset classifies people described by a set of attributes as
good or bad credit risks. Comes in two formats (one all numeric). Also comes
with a cost matrix', 'area': 'Social Science', 'tasks': ['Classification'],
'characteristics': ['Multivariate'], 'num_instances': 1000, 'num_features': 20,
'feature_types': ['Categorical', 'Integer'], 'demographics': ['Other', 'Marital
Status', 'Age', 'Occupation'], 'target_col': ['class'], 'index_col': None,
'has_missing_values': 'no', 'missing_values_symbol': None,
'year_of_dataset_creation': 1994, 'last_updated': 'Thu Aug 10 2023',
'dataset_doi': '10.24432/C5NC77', 'creators': ['Hans Hofmann'], 'intro_paper':
None, 'additional_info': {'summary': 'Two datasets are provided. the original
dataset, in the form provided by Prof. Hofmann, contains categorical/symbolic
attributes and is in the file "german.data". \r\n\r\nFor algorithms that need
numerical attributes, Strathclyde University produced the file "german.data-
numeric". This file has been edited and several indicator variables added to
make it suitable for algorithms which cannot cope with categorical variables.
Several attributes that are ordered categorical (such as attribute 17) have been
coded as integer. This was the form used by StatLog.\r\n\r\nThis dataset
requires use of a cost matrix (see below)\r\n\r\n ... 1
2\r\n-----\r\n 1  0
1\r\n-----\r\n 2  5      0\r\n\r\n(1 = Good, 2 =
Bad)\r\n\r\nThe rows represent the actual classification and the columns the
predicted classification.\r\n\r\nIt is worse to class a customer as good when
they are bad (5), than it is to class a customer as bad when they are good
(1).\r\n', 'purpose': None, 'funded_by': None, 'instances_represent': None,
'recommended_data_splits': None, 'sensitive_data': None,
'preprocessing_description': None, 'variable_info': 'Attribute 1: (qualitative)
\r\n Status of existing checking account\r\n          A11 :      ... <    0
DM\r\n\t          A12 : 0 <= ... < 200 DM\r\n\t          A13 :      ... >= 200 DM /
salary assignments for at least 1 year\r\n          A14 : no checking
account\r\n\r\nAttribute 2: (numerical)\r\n\t          Duration in
month\r\n\r\nAttribute 3: (qualitative)\r\n\t          Credit history\r\n\t
```

A30 : no credits taken/ all credits paid back duly\r\n A31 : all
credits at this bank paid back duly\r\n\t A32 : existing credits paid back
duly till now\r\n A33 : delay in paying off in the past\r\n\t
A34 : critical account/ other credits existing (not at this
bank)\r\n\r\nAttribute 4: (qualitative)\r\n\t Purpose\r\n\t A40 : car
(new)\r\n\t A41 : car (used)\r\n\t A42 : furniture/equipment\r\n\t
A43 : radio/television\r\n\t A44 : domestic appliances\r\n\t A45 :
repairs\r\n\t A46 : education\r\n\t A47 : (vacation - does not
exist?)\r\n\t A48 : retraining\r\n\t A49 : business\r\n\t A410 :
others\r\n\r\nAttribute 5: (numerical)\r\n\t Credit amount\r\n\r\nAttribute
6: (qualitative)\r\n\t Savings account/bonds\r\n\t A61 : ...
< 100 DM\r\n\t A62 : 100 <= ... < 500 DM\r\n\t A63 : 500 <= ...
< 1000 DM\r\n\t A64 : .. >= 1000 DM\r\n A65 :
unknown/ no savings account\r\n\r\nAttribute 7: (qualitative)\r\n\t
Present employment since\r\n\t A71 : unemployed\r\n\t A72 : ...
< 1 year\r\n\t A73 : 1 <= ... < 4 years \r\n\t A74 : 4 <= ... < 7
years\r\n\t A75 : .. >= 7 years\r\n\r\nAttribute 8:
(numerical)\r\n\t Installment rate in percentage of disposable
income\r\n\r\nAttribute 9: (qualitative)\r\n\t Personal status and
sex\r\n\t A91 : male : divorced/separated\r\n\t A92 : female :
divorced/separated/married\r\n A93 : male : single\r\n\t A94
: male : married/widowed\r\n\t A95 : female : single\r\n\r\nAttribute 10:
(qualitative)\r\n\t Other debtors / guarantors\r\n\t A101 : none\r\n\t
A102 : co-applicant\r\n\t A103 : guarantor\r\n\r\nAttribute 11:
(numerical)\r\n\t Present residence since\r\n\r\nAttribute 12:
(qualitative)\r\n\t Property\r\n\t A121 : real estate\r\n\t A122
: if not A121 : building society savings agreement/ life insurance\r\n
A123 : if not A121/A122 : car or other, not in attribute 6\r\n\t A124 :
unknown / no property\r\n\r\nAttribute 13: (numerical)\r\n\t Age in
years\r\n\r\nAttribute 14: (qualitative)\r\n\t Other installment plans
\r\n\t A141 : bank\r\n\t A142 : stores\r\n\t A143 :
none\r\n\r\nAttribute 15: (qualitative)\r\n\t Housing\r\n\t A151 :
rent\r\n\t A152 : own\r\n\t A153 : for free\r\n\r\nAttribute 16:
(numerical)\r\n Number of existing credits at this
bank\r\n\r\nAttribute 17: (qualitative)\r\n\t Job\r\n\t A171 :
unemployed/ unskilled - non-resident\r\n\t A172 : unskilled -
resident\r\n\t A173 : skilled employee / official\r\n\t A174 :
management/ self-employed/\r\n\t\t highly qualified employee/
officer\r\n\r\nAttribute 18: (numerical)\r\n\t Number of people being
liable to provide maintenance for\r\n\r\nAttribute 19: (qualitative)\r\n\t
Telephone\r\n\t A191 : none\r\n\t A192 : yes, registered under the
customers name\r\n\r\nAttribute 20: (qualitative)\r\n\t foreign
worker\r\n\t A201 : yes\r\n\t A202 : no\r\n\t, 'citation': None}}

	name	role	type	demographic \
0	Attribute1	Feature	Categorical	None
1	Attribute2	Feature	Integer	None
2	Attribute3	Feature	Categorical	None
3	Attribute4	Feature	Categorical	None

4	Attribute5	Feature	Integer	None
5	Attribute6	Feature	Categorical	None
6	Attribute7	Feature	Categorical	Other
7	Attribute8	Feature	Integer	None
8	Attribute9	Feature	Categorical	Marital Status
9	Attribute10	Feature	Categorical	None
10	Attribute11	Feature	Integer	None
11	Attribute12	Feature	Categorical	None
12	Attribute13	Feature	Integer	Age
13	Attribute14	Feature	Categorical	None
14	Attribute15	Feature	Categorical	Other
15	Attribute16	Feature	Integer	None
16	Attribute17	Feature	Categorical	Occupation
17	Attribute18	Feature	Integer	None
18	Attribute19	Feature	Binary	None
19	Attribute20	Feature	Binary	Other
20	class	Target	Binary	None

		description	units	missing_values
0		Status of existing checking account	None	no
1		Duration	months	no
2		Credit history	None	no
3		Purpose	None	no
4		Credit amount	None	no
5		Savings account/bonds	None	no
6		Present employment since	None	no
7	Installment rate in percentage of disposable i...		None	no
8		Personal status and sex	None	no
9		Other debtors / guarantors	None	no
10		Present residence since	None	no
11		Property	None	no
12		Age	years	no
13		Other installment plans	None	no
14		Housing	None	no
15	Number of existing credits at this bank		None	no
16		Job	None	no
17	Number of people being liable to provide maint...		None	no
18		Telephone	None	no
19		foreign worker	None	no
20		1 = Good, 2 = Bad	None	no

```
[108]: from sklearn.preprocessing import LabelEncoder, LabelBinarizer
```

```
statlog_german_credit_data_df = pd.DataFrame(
    ↪from_dict(statlog_german_credit_data.data.features)
statlog_german_credit_data_df['class'] = statlog_german_credit_data.data.
    ↪targets['class']
```

```

statlog_german_credit_data_df['class'] = statlog_german_credit_data_df['class'].
↳ astype(str)

encoded_features = statlog_german_credit_data_df.columns
label_encoders = {feature: LabelEncoder() for feature in encoded_features}
for feature, encoder in label_encoders.items():
    statlog_german_credit_data_df[feature] = encoder.
↳ fit_transform(statlog_german_credit_data_df[feature])

statlog_german_credit_data_df

```

```

[108]:
Attribute1 Attribute2 Attribute3 Attribute4 Attribute5 Attribute6 \
0          0          2          4          4         142          4
1          1         29          2          4         770          0
2          3          8          4          7         390          0
3          0         26          2          3         848          0
4          0         17          3          0         734          0
..         ...         ...         ...         ...         ...         ...
995         3          8          2          3         310          0
996         0         21          2          1         661          0
997         3          8          2          4          69          0
998         0         27          2          4         332          0
999         1         27          4          1         711          1

```

```

Attribute7 Attribute8 Attribute9 Attribute10 ... Attribute12 \
0          4          3          2          0 ...          0
1          2          1          1          0 ...          0
2          3          1          2          0 ...          0
3          3          1          2          2 ...          1
4          2          2          2          0 ...          3
..         ...         ...         ...         ...         ...
995         3          2          1          0 ...          0
996         2          3          0          0 ...          1
997         4          3          2          0 ...          2
998         2          3          2          0 ...          3
999         0          2          2          0 ...          2

```

```

Attribute13 Attribute14 Attribute15 Attribute16 Attribute17 \
0          48          2          1          1          2
1          3          2          1          0          2
2          30          2          1          0          1
3          26          2          2          0          2
4          34          2          2          1          2
..         ...         ...         ...         ...         ...
995         12          2          1          0          1
996         21          2          1          0          3
997         19          2          1          0          2

```

998	4	2	2	0	2
999	8	2	1	0	2

	Attribute18	Attribute19	Attribute20	class
0	0	1	0	0
1	0	0	0	1
2	1	0	0	0
3	1	0	0	0
4	1	0	0	1
..
995	0	0	0	0
996	0	1	0	0
997	0	0	0	0
998	0	1	0	1
999	0	0	0	0

[1000 rows x 21 columns]

```
[110]: X = statlog_german_credit_data_df.drop(['class'], axis='columns')
y = statlog_german_credit_data_df['class']
```

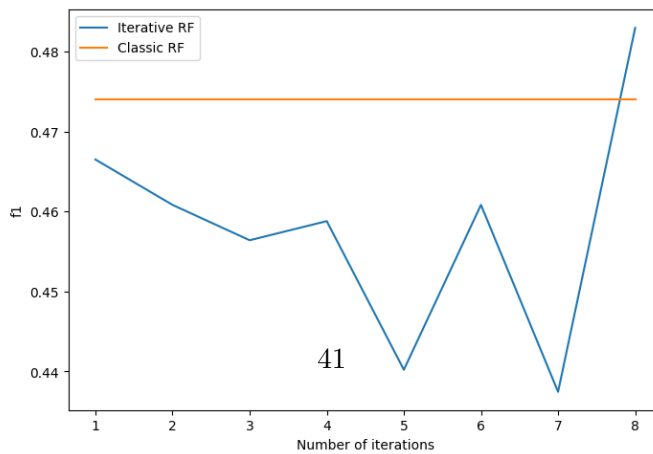
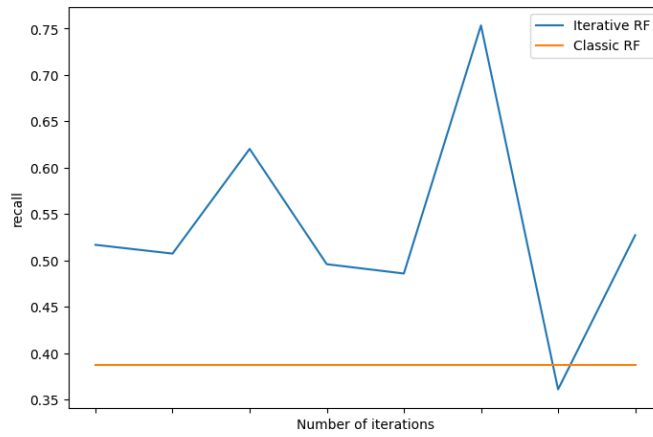
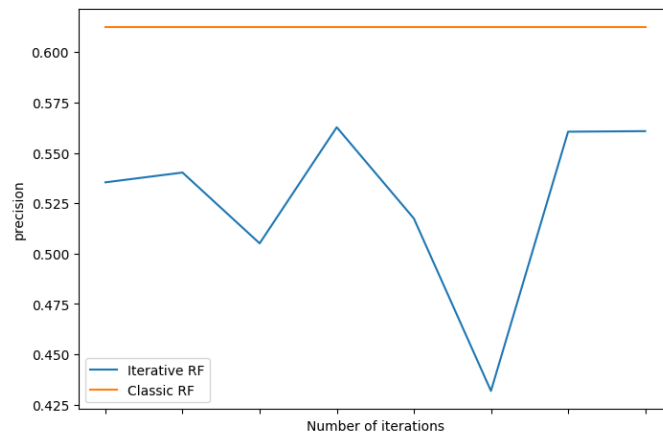
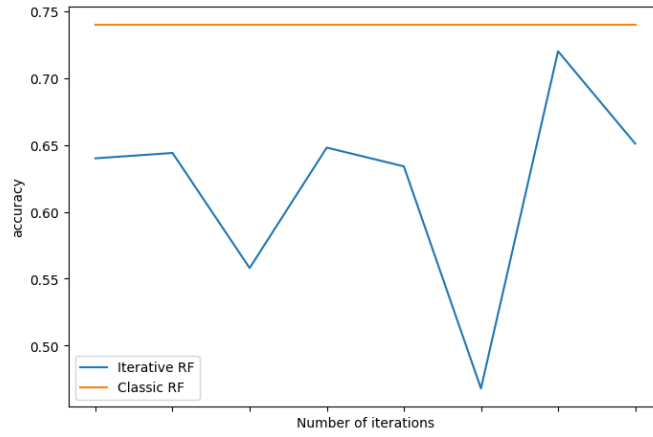
```
[120]: NUM_ESTIMATORS = 10
MAX_DEPTH = 15
iteration_counts = list(range(1, 9))
measures = ['accuracy', 'precision', 'recall', 'f1']
res = [evaluate(IterativeRandomForrest(n_estimators=NUM_ESTIMATORS,
    ↳max_depth=MAX_DEPTH, iter_count=iter_count), X, y, scoring=measures) for
    ↳iter_count in iteration_counts]

res_ref_single = evaluate(RandomForestClassifier(
    max_depth=MAX_DEPTH,
    n_estimators=NUM_ESTIMATORS,
    n_jobs=8
), X, y, scoring=measures)
res_ref = [res_ref_single] * len(iteration_counts)

fig, axes = plt.subplots(4, 1, sharex=True)
fig.set_size_inches(8, 25)
for measure, ax in zip(measures, axes):
    _y_iter = [r[f"test_{measure}"] for r in res]
    _y_ref = [r[f"test_{measure}"] for r in res_ref]
    ax.plot(iteration_counts, _y_iter, label="Iterative RF")
    ax.plot(iteration_counts, _y_ref, label="Classic RF")
    ax.set_xlabel("Number of iterations")
```

```
ax.set_ylabel(measure)
ax.legend()

plt.show()
```

```
[171]: class IterativeRandomForrestSampleBased(BaseEstimator):
    def __init__(
        self,
        *,
        iter_count=2,
        n_estimators=2,
        max_depth=None,
    ):

        self.iter_count = iter_count
        self.n_estimators = n_estimators
        self.max_depth = max_depth
        self._fitted_estimator = None

    def _get_estimator(self):
        return RandomForestClassifier(
            n_estimators=self.n_estimators,
            max_depth=self.max_depth,
            n_jobs=8
        )

    def fit(self, X, y):
        for iteration_idx in range(self.iter_count):
            estimator = self._get_estimator()
            estimator.fit(X, y)
            y_pred = estimator.predict(X)
            misclassified_samples = np.where(y_pred != y)[0]
            X = pd.concat([X, X.iloc[misclassified_samples]])
            y = np.concatenate([y, y[misclassified_samples]])

            self._fitted_estimator = estimator

    def get_fiited_estimator(self):
        return self._fitted_estimator

    def predict(self, X):
        return self._fitted_estimator.predict(X)

    def decision_function(self, X):
        # called for roc_auc_score
        raise NotImplementedError("This IterativeRandomForrest instance is not_
↪fitted yet")
```

```

def predict_proba(self, X):
    # called for roc_auc_score
    raise NotImplementedError("This IterativeRandomForrest instance is not_
↳fitted yet")

clf = IterativeRandomForrestSampleBased(iter_count=2, n_estimators=10,
↳max_depth=15)
clf.fit(X, y)
fancy_rf = clf.get_fiited_estimator()
evaluate(fancy_rf, X, y, scoring=['accuracy', 'precision', 'recall', 'f1'])
# evaluate(clf, X, y, scoring=['accuracy', 'precision', 'recall', 'f1'])

```

```

[171]: {'fit_time': 0.03116297721862793,
'score_time': 0.021782112121582032,
'test_accuracy': 0.734,
'train_accuracy': 0.9807499999999999,
'test_precision': 0.6241318896157606,
'train_precision': 0.9965100519539192,
'test_recall': 0.33318823460535785,
'train_recall': 0.9391317615112357,
'test_f1': 0.4290354054880171,
'train_f1': 0.96685846006811}

```

```
[ ]:
```

```

[173]: NUM_ESTIMATORS = 2
MAX_DEPTH = 2
iteration_counts = list(range(1, 9))
measures = ['accuracy', 'precision', 'recall', 'f1']

itertative_rfs = []
for iter_count in iteration_counts:
    clf = IterativeRandomForrestSampleBased(n_estimators=NUM_ESTIMATORS,
↳max_depth=MAX_DEPTH, iter_count=iter_count)
    clf.fit(X, y)
    itertative_rfs.append(clf.get_fiited_estimator())

res = [evaluate(itertative_rfs[idx], X, y, scoring=measures) for idx,
↳iter_count in enumerate(iteration_counts)]

res_ref_single = evaluate(RandomForestClassifier(
    max_depth=MAX_DEPTH,
    n_estimators=NUM_ESTIMATORS,
    n_jobs=8
), X, y, scoring=measures)
res_ref = [res_ref_single] * len(iteration_counts)

```

```

fig, axes = plt.subplots(4, 1, sharex=True)
fig.set_size_inches(8, 25)
for measure, ax in zip(measures, axes):
    _y_iter = [r[f"test_{measure}"] for r in res]
    _y_ref = [r[f"test_{measure}"] for r in res_ref]
    ax.plot(iteration_counts, _y_iter, label="Iterative RF")
    ax.plot(iteration_counts, _y_ref, label="Classic RF")
    ax.set_xlabel("Number of iterations")
    ax.set_ylabel(measure)
    ax.legend()

plt.show()

```

```

/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/metrics/_classification.py:1497: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 due to no predicted samples. Use
`zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/metrics/_classification.py:1497: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 due to no predicted samples. Use
`zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/metrics/_classification.py:1497: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 due to no predicted samples. Use
`zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/metrics/_classification.py:1497: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 due to no predicted samples. Use
`zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/metrics/_classification.py:1497: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 due to no predicted samples. Use
`zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/metrics/_classification.py:1497: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 due to no predicted samples. Use
`zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

```

`zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/metrics/_classification.py:1497: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 due to no predicted samples. Use
`zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/metrics/_classification.py:1497: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 due to no predicted samples. Use
`zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/metrics/_classification.py:1497: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 due to no predicted samples. Use
`zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/media/data/coding/uma/.venv/lib/python3.11/site-
packages/sklearn/metrics/_classification.py:1497: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 due to no predicted samples. Use
`zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

